

A_12 (1)

June 1, 2019

```
In [0]: import matplotlib.pyplot as plt
        from keras.utils import np_utils
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.layers import Dropout
        from keras.layers.normalization import BatchNormalization
        from keras.optimizers import Adam
        from keras.layers import Dense, Activation
        import seaborn as sns
        import numpy as np
        import keras
```

```
%matplotlib inline
```

```
In [2]: (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
In [3]: print("Training data shape: ", X_train.shape) # (60000, 28, 28) -- 60000 images, each 28x28
        print("Test data shape", X_test.shape)
```

```
        print("Training label shape: ", y_train.shape) # (60000,) -- 60000 numbers (all 0-9)
        print("First 5 training labels: ", y_train[:5]) # [5, 0, 4, 1, 9]
```

Training data shape: (60000, 28, 28)

Test data shape (10000, 28, 28)

Training label shape: (60000,)

First 5 training labels: [5 0 4 1 9]

```
In [0]: # Flatten the images
        image_vector_size = 28*28
        X_train = X_train.reshape(X_train.shape[0], image_vector_size)
        X_test = X_test.reshape(X_test.shape[0], image_vector_size)
```

```
In [5]: # Convert to "one-hot" vectors using the to_categorical function
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
print("First 5 training lables as one-hot encoded vectors:\n", y_train[:5])
```

First 5 training lables as one-hot encoded vectors:

```
[[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]]
```

0.1 Model 1 without dropout and BN

```
In [0]: output_dim = 10
input_dim = X_train.shape[1]
```

```
nb_epoch = 10
batch_size = 128
```

```
In [0]: model1 = Sequential()
```

```
model1.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer='random_uniform'))
model1.add(Dense(120, activation='relu', kernel_initializer='random_uniform'))
```

```
model1.add(Dense(output_dim, activation='softmax'))
```

```
print(model1.summary())
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_15 (Dense)             (None, 364)              285740
-----
dense_16 (Dense)             (None, 120)              43800
-----
dense_17 (Dense)             (None, 10)               1210
=====
Total params: 330,750
Trainable params: 330,750
Non-trainable params: 0
-----
None
```

```
In [0]: model1.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```

In [0]: # Training the model
        history = model1.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20


In [0]: score = model1.evaluate(X_test, y_test)

10000/10000 [=====] - 0s 49us/step



In [0]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();

```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.8604

0.1.1 2 Layers with BN and Dropout

```
In [0]: model2 = Sequential()
```

```
model2.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
```

```
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
```

```
model2.add(Dense(output_dim, activation='softmax'))
```

```
print(model2.summary())
```

Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 364)	285740
dense_22 (Dense)	(None, 128)	46720

```

-----
batch_normalization_3 (Batch Normalization)      512
-----
dropout_3 (Dropout)                             0
-----
dense_23 (Dense)                                1290
=====
Total params: 334,262
Trainable params: 334,006
Non-trainable params: 256
-----
None

```

```

In [0]: model2.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```

In [0]: # Training the model
        history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

```

```

In [0]: score = model2.evaluate(X_test, y_test)

```

```

10000/10000 [=====] - 1s 55us/step

```

```


In [0]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();

```



```

In [0]: print('Test accuracy:', score[1])

```

Test accuracy: 0.9819

0.1.2 Changing dropout rate to 0.8

```

In [16]: model2 = Sequential()

        model2.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer=

```

```

model2.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))

model2.add(BatchNormalization())
model2.add(Dropout(0.8))

model2.add(Dense(output_dim, activation='softmax'))

print(model2.summary())

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:307: Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 364)	285740
dense_2 (Dense)	(None, 128)	46720
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290

Total params: 334,262
 Trainable params: 334,006
 Non-trainable params: 256

```

In [0]: model2.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])

```

```

In [19]: # Training the model
         history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1)

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1610 - acc: 0.9561 - val_loss: 0.1238
Epoch 2/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1238 - acc: 0.9661 - val_loss: 0.1043
Epoch 3/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1043 - acc: 0.9710 - val_loss: 0.0880
Epoch 4/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0880 - acc: 0.9758 - val_loss: 0.0880
Epoch 5/20

```

```

60000/60000 [=====] - 6s 93us/step - loss: 0.0741 - acc: 0.9792 - val.
Epoch 6/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0638 - acc: 0.9816 - val.
Epoch 7/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0622 - acc: 0.9829 - val.
Epoch 8/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0528 - acc: 0.9850 - val.
Epoch 9/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0472 - acc: 0.9863 - val.
Epoch 10/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0426 - acc: 0.9875 - val.
Epoch 11/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0409 - acc: 0.9878 - val.
Epoch 12/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0345 - acc: 0.9894 - val.
Epoch 13/20
60000/60000 [=====] - 5s 92us/step - loss: 0.0312 - acc: 0.9905 - val.
Epoch 14/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0307 - acc: 0.9911 - val.
Epoch 15/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0285 - acc: 0.9915 - val.
Epoch 16/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0288 - acc: 0.9913 - val.
Epoch 17/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0265 - acc: 0.9922 - val.
Epoch 18/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0237 - acc: 0.9927 - val.
Epoch 19/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0234 - acc: 0.9930 - val.
Epoch 20/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0185 - acc: 0.9945 - val.

```

```
In [20]: score = model2.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 0s 50us/step
```

```

In [21]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))


        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")

```



```
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [22]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9803

0.1.3 Changing dropout rate to 0.2

```
In [30]: model2 = Sequential()
```

```
model2.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer=
model2.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
```

```
model2.add(BatchNormalization())
model2.add(Dropout(0.2))
```

```
model2.add(Dense(output_dim, activation='softmax'))
```

```
print(model2.summary())
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
dense_7 (Dense)                (None, 364)                285740
-----
dense_8 (Dense)                (None, 128)                46720
-----
batch_normalization_3 (Batch (None, 128)                512
-----
dropout_3 (Dropout)           (None, 128)                0
-----
dense_9 (Dense)                (None, 10)                 1290
=====
Total params: 334,262
Trainable params: 334,006
Non-trainable params: 256
-----
None

```

```

In [0]: model2.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```

In [32]: # Training the model

```

```

        history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/10
60000/60000 [=====] - 6s 101us/step - loss: 0.2233 - acc: 0.9340 - val.
Epoch 2/10
60000/60000 [=====] - 5s 89us/step - loss: 0.0950 - acc: 0.9714 - val.
Epoch 3/10
60000/60000 [=====] - 5s 89us/step - loss: 0.0646 - acc: 0.9803 - val.
Epoch 4/10
60000/60000 [=====] - 5s 90us/step - loss: 0.0496 - acc: 0.9841 - val.
Epoch 5/10
60000/60000 [=====] - 5s 89us/step - loss: 0.0382 - acc: 0.9874 - val.
Epoch 6/10
60000/60000 [=====] - 5s 90us/step - loss: 0.0315 - acc: 0.9902 - val.
Epoch 7/10
60000/60000 [=====] - 5s 90us/step - loss: 0.0269 - acc: 0.9912 - val.
Epoch 8/10
60000/60000 [=====] - 6s 92us/step - loss: 0.0237 - acc: 0.9925 - val.
Epoch 9/10
60000/60000 [=====] - 6s 92us/step - loss: 0.0215 - acc: 0.9926 - val.
Epoch 10/10
60000/60000 [=====] - 5s 90us/step - loss: 0.0180 - acc: 0.9941 - val.

```

```

In [33]: score = model2.evaluate(X_test, y_test)

```


10000/10000 [=====] - 0s 48us/step

```
In [34]: fig,ax = plt.subplots(1,1)
         ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

         # list of epoch numbers
         list_of_epoch = list(range(1,nb_epoch+1))

         train_loss = history.history['loss']
         val_loss = history.history['val_loss']

         ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
         ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
         plt.legend()
         plt.grid()
         plt.show();
```



```
In [35]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9813

0.14 Model with 3 hidden layers

```
In [0]: model3 = Sequential()
```

```

model3.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer='random_uniform'))
model3.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
model3.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))

#model2.add(BatchNormalization())
#model2.add(Dropout(0.5))

model3.add(Dense(output_dim, activation='softmax'))

print(model3.summary())

```

```

-----
Layer (type)                 Output Shape              Param #
-----
dense_28 (Dense)             (None, 364)              285740
-----
dense_29 (Dense)             (None, 128)              46720
-----
dense_30 (Dense)             (None, 64)               8256
-----
dense_31 (Dense)             (None, 10)               650
=====
Total params: 341,366
Trainable params: 341,366
Non-trainable params: 0
-----
None

```

```

In [0]: model3.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```

In [0]: # Training the model
        history = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20

```

Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20


```
In [0]: score = model3.evaluate(X_test, y_test)
10000/10000 [=====] - 1s 53us/step
```

```
In [0]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9774

0.1.5 Model with 3 hidden layers along with BN and dropout

```
In [0]: model4 = Sequential()
```

```
model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model4.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
```

```
model4.add(BatchNormalization())
model4.add(Dropout(0.5))
```

```
model4.add(Dense(output_dim, activation='softmax'))
```

```
print(model4.summary())
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_32 (Dense)             (None, 364)              285740
-----
dense_33 (Dense)             (None, 128)              46720
-----
dense_34 (Dense)             (None, 64)               8256
-----
batch_normalization_4 (Batch Normalization) (None, 64)              256
-----
dropout_4 (Dropout)          (None, 64)               0
-----
dense_35 (Dense)             (None, 10)               650
=====
Total params: 341,622
Trainable params: 341,494
Non-trainable params: 128
-----
None
```

```
In [0]: model4.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```

In [0]: # Training the model
        history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

In [0]: score = model4.evaluate(X_test, y_test)

10000/10000 [=====] - 1s 58us/step


In [0]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();

```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9837

0.1.6 changing dropout rate to 0.8

```
In [36]: model4 = Sequential()
```

```
model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=
model4.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
```

```
model4.add(BatchNormalization())
model4.add(Dropout(0.8))
```

```
model4.add(Dense(output_dim, activation='softmax'))
```

```
print(model4.summary())
```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 364)	285740

dense_11 (Dense)	(None, 128)	46720

dense_12 (Dense)	(None, 64)	8256

batch_normalization_4 (Batch Normalization)	(None, 64)	256

dropout_4 (Dropout)	(None, 64)	0

dense_13 (Dense)	(None, 10)	650
=====		
Total params: 341,622		
Trainable params: 341,494		
Non-trainable params: 128		

None		

```
In [0]: model4.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [38]: # Training the model
```

```
        history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/10
60000/60000 [=====] - 6s 104us/step - loss: 0.7021 - acc: 0.7874 - val_loss: 0.3500
Epoch 2/10
60000/60000 [=====] - 5s 90us/step - loss: 0.2923 - acc: 0.9194 - val_loss: 0.1800
Epoch 3/10
60000/60000 [=====] - 5s 91us/step - loss: 0.1997 - acc: 0.9452 - val_loss: 0.1200
Epoch 4/10
60000/60000 [=====] - 5s 91us/step - loss: 0.1612 - acc: 0.9552 - val_loss: 0.0900
Epoch 5/10
60000/60000 [=====] - 5s 92us/step - loss: 0.1353 - acc: 0.9605 - val_loss: 0.0700
Epoch 6/10
60000/60000 [=====] - 5s 90us/step - loss: 0.1188 - acc: 0.9655 - val_loss: 0.0500
Epoch 7/10
60000/60000 [=====] - 5s 91us/step - loss: 0.1013 - acc: 0.9692 - val_loss: 0.0400
Epoch 8/10
60000/60000 [=====] - 5s 89us/step - loss: 0.0933 - acc: 0.9719 - val_loss: 0.0300
Epoch 9/10
60000/60000 [=====] - 5s 90us/step - loss: 0.0830 - acc: 0.9746 - val_loss: 0.0200
Epoch 10/10
60000/60000 [=====] - 5s 88us/step - loss: 0.0747 - acc: 0.9760 - val_loss: 0.0100
```

```
In [39]: score = model4.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 51us/step

```
In [40]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();
```



```
In [41]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9767

0.1.7 changing dropout rate to 0.2

```
In [42]: model4 = Sequential()
```

```

model4.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer=
model4.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))

model4.add(BatchNormalization())
model4.add(Dropout(0.2))

model4.add(Dense(output_dim, activation='softmax'))

print(model4.summary())

```

```

-----
Layer (type)                Output Shape                Param #
-----
dense_14 (Dense)            (None, 364)                285740
-----
dense_15 (Dense)            (None, 128)                46720
-----
dense_16 (Dense)            (None, 64)                 8256
-----
batch_normalization_5 (Batch Normalization) (None, 64)                256
-----
dropout_5 (Dropout)         (None, 64)                 0
-----
dense_17 (Dense)            (None, 10)                 650
-----
Total params: 341,622
Trainable params: 341,494
Non-trainable params: 128
-----
None

```

```

In [0]: model4.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```

In [44]: # Training the model

```

```

        history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 105us/step - loss: 0.2361 - acc: 0.9319 - val_loss: 0.1051 - val_acc: 0.9717

Epoch 2/10

60000/60000 [=====] - 5s 89us/step - loss: 0.0951 - acc: 0.9717 - val_loss: 0.0647 - val_acc: 0.9810

Epoch 3/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0647 - acc: 0.9810 - val_loss: 0.0511 - val_acc: 0.9843

Epoch 4/10

60000/60000 [=====] - 5s 88us/step - loss: 0.0511 - acc: 0.9843 - val_loss: 0.0478 - val_acc: 0.9875

```

Epoch 5/10
60000/60000 [=====] - 5s 88us/step - loss: 0.0386 - acc: 0.9879 - val.
Epoch 6/10
60000/60000 [=====] - 5s 90us/step - loss: 0.0332 - acc: 0.9896 - val.
Epoch 7/10
60000/60000 [=====] - 5s 87us/step - loss: 0.0248 - acc: 0.9918 - val.
Epoch 8/10
60000/60000 [=====] - 5s 88us/step - loss: 0.0235 - acc: 0.9921 - val.
Epoch 9/10
60000/60000 [=====] - 5s 87us/step - loss: 0.0196 - acc: 0.9939 - val.
Epoch 10/10
60000/60000 [=====] - 5s 88us/step - loss: 0.0195 - acc: 0.9937 - val.

```

```
In [45]: score = model4.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 1s 55us/step
```

```


In [46]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();

```



```
In [47]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9813

0.1.8 Model with 5 hidden layers

```
In [0]: model5 = Sequential()
```

```
model5.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model5.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
```

```
#model2.add(BatchNormalization())
#model2.add(Dropout(0.5))
```

```
model5.add(Dense(output_dim, activation='softmax'))
```

```
print(model5.summary())
```

Layer (type)	Output Shape	Param #
=====		

dense_36 (Dense)	(None, 364)	285740

dense_37 (Dense)	(None, 128)	46720

dense_38 (Dense)	(None, 96)	12384

dense_39 (Dense)	(None, 64)	6208

dense_40 (Dense)	(None, 32)	2080

dense_41 (Dense)	(None, 10)	330
=====		
Total params: 353,462		
Trainable params: 353,462		
Non-trainable params: 0		

None		

```
In [0]: model5.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
        history = model5.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

```
In [0]: score = model15.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 61us/step

```
In [0]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9804

0.1.9 Model with 5 hidden layers along with BN and Dropout

```
In [0]: model6 = Sequential()
```

```
model6.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
```

```
model6.add(BatchNormalization())
model6.add(Dropout(0.5))
```

```
model6.add(Dense(output_dim, activation='softmax'))
```

```
print(model6.summary())
```

Layer (type)	Output Shape	Param #
dense_48 (Dense)	(None, 364)	285740
dense_49 (Dense)	(None, 128)	46720
dense_50 (Dense)	(None, 96)	12384
dense_51 (Dense)	(None, 64)	6208
dense_52 (Dense)	(None, 32)	2080
batch_normalization_6 (Batch Normalization)	(None, 32)	128
dropout_6 (Dropout)	(None, 32)	0
dense_53 (Dense)	(None, 10)	330

=====
Total params: 353,590
Trainable params: 353,526
Non-trainable params: 64
=====
None

```
In [0]: model6.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
```

```
history = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

```
In [0]: score = model6.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 83us/step

```
In [0]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9814

0.1.10 Changing Dropout to 0.8

```
In [48]: model6 = Sequential()
```

```
model6.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer=
model6.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))

model6.add(BatchNormalization())
model6.add(Dropout(0.8))

model6.add(Dense(output_dim, activation='softmax'))

print(model6.summary())
```

```
-----
Layer (type)                Output Shape                Param #
=====
```

dense_18 (Dense)	(None, 364)	285740

dense_19 (Dense)	(None, 128)	46720

dense_20 (Dense)	(None, 96)	12384

dense_21 (Dense)	(None, 64)	6208

dense_22 (Dense)	(None, 32)	2080

batch_normalization_6 (Batch Normalization)	(None, 32)	128

dropout_6 (Dropout)	(None, 32)	0

dense_23 (Dense)	(None, 10)	330
=====		
Total params: 353,590		
Trainable params: 353,526		
Non-trainable params: 64		

None		

```
In [0]: model6.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [50]: # Training the model
```

```
        history = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

Epoch 2/10

Epoch 3/10

Epoch 4/10

Epoch 5/10

Epoch 6/10

Epoch 7/10

Epoch 8/10

Epoch 9/10

Epoch 10/10

```
In [51]: score = model6.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 1s 56us/step
```

```
In [52]: fig,ax = plt.subplots(1,1)
```

```
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')
```


```

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```



```
In [53]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9774

0.1.11 Changing dropout to 0.2

```
In [54]: model6 = Sequential()
```

```

model6.add(Dense(364, activation='relu', input_shape=(input_dim,),kernel_initializer=
model6.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))

```

```

model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))

model6.add(BatchNormalization())
model6.add(Dropout(0.2))

model6.add(Dense(output_dim, activation='softmax'))

print(model6.summary())

```

```

-----
Layer (type)                 Output Shape              Param #
=====
dense_24 (Dense)             (None, 364)              285740
-----
dense_25 (Dense)             (None, 128)              46720
-----
dense_26 (Dense)             (None, 96)               12384
-----
dense_27 (Dense)             (None, 64)               6208
-----
dense_28 (Dense)             (None, 32)               2080
-----
batch_normalization_7 (Batch Normalization) (None, 32)              128
-----
dropout_7 (Dropout)          (None, 32)               0
-----
dense_29 (Dense)             (None, 10)               330
=====
Total params: 353,590
Trainable params: 353,526
Non-trainable params: 64
-----
None

```

```

In [0]: model6.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```

In [57]: # Training the model

```

```

        history = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=0)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 96us/step - loss: 0.1093 - acc: 0.9702 - val_loss: 0.0770 - val_acc: 0.9781

Epoch 2/10

60000/60000 [=====] - 6s 94us/step - loss: 0.0770 - acc: 0.9781 - val_loss: 0.0770 - val_acc: 0.9781

Epoch 3/10

```

60000/60000 [=====] - 6s 96us/step - loss: 0.0563 - acc: 0.9838 - val.
Epoch 4/10
60000/60000 [=====] - 6s 97us/step - loss: 0.0462 - acc: 0.9865 - val.
Epoch 5/10
60000/60000 [=====] - 6s 95us/step - loss: 0.0383 - acc: 0.9890 - val.
Epoch 6/10
60000/60000 [=====] - 6s 95us/step - loss: 0.0338 - acc: 0.9899 - val.
Epoch 7/10
60000/60000 [=====] - 6s 93us/step - loss: 0.0286 - acc: 0.9916 - val.
Epoch 8/10
60000/60000 [=====] - 6s 94us/step - loss: 0.0280 - acc: 0.9916 - val.
Epoch 9/10
60000/60000 [=====] - 6s 94us/step - loss: 0.0242 - acc: 0.9927 - val.
Epoch 10/10
60000/60000 [=====] - 6s 95us/step - loss: 0.0213 - acc: 0.9934 - val.

```

```
In [59]: score = model6.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 1s 65us/step
```


```
In [60]: fig,ax = plt.subplots(1,1)
        ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

        # list of epoch numbers
        list_of_epoch = list(range(1,nb_epoch+1))

        train_loss = history.history['loss']
        val_loss = history.history['val_loss']

        ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
        ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
        plt.legend()
        plt.grid()
        plt.show();

```



```
In [58]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9774

1 Conclusions

```
In [64]: # Please compare all your models using Prettytable library
         from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Númer of Layers", "BN", "Dropout", "Accuracy"]
```

```
x.add_row(["2", 'NO', "NO", 0.86])
```

```
x.add_row(["2", 'YES', 0.5, 0.98])
```

```
x.add_row(["2", 'YES', 0.8, 0.98])
```

```
x.add_row(["2", 'YES', 0.2, 0.98])
```

```
x.add_row(["3", 'NO', "NO", 0.87])
```

```
x.add_row(["3", 'YES', 0.5, 0.98])
```

```
x.add_row(["3", 'YES', 0.8, 0.97])
```

```
x.add_row(["3", 'YES', 0.2, 0.98])
```

```
x.add_row(["5", 'NO', 'NO', 0.98])
```

```
x.add_row(["5", 'YES', 0.5, 0.98])
```

```
x.add_row(["5", 'YES', 0.8, 0.97])
```

```
x.add_row(["5", 'YES', 0.2, 0.97])
```

```
print(x)
```

Numer of Layers	BN	Dropout	Accuracy
2	NO	NO	0.86
2	YES	0.5	0.98
2	YES	0.8	0.98
2	YES	0.2	0.98
3	NO	NO	0.87
3	YES	0.5	0.98
3	YES	0.8	0.97
3	YES	0.2	0.98
5	NO	NO	0.98
5	YES	0.5	0.98
5	YES	0.8	0.97
5	YES	0.2	0.97